

Lecture 10.1 : Recursion

Introduction ¶

- A *recursive* solution is one where the solution to a problem is expressed as an operation on a *simplified* version of the *same* problem.
- For certain problems, recursion may offer an intuitive, simple, and elegant solution. The ability to both recognise a problem that lends itself to a recursive solution and to implement that solution is an important skill that will make you a better programmer. Furthermore, some programming languages, such as Prolog (which you will meet in second year), make heavy use of recursion.
- We introduce recursion below and implement, in Python, recursive solutions to a selection of programming problems.

What is recursion?

- Any function that calls itself is *recursive* and exhibits *recursion*.

```
def foo(n):  
    return foo(n-1)
```

- The function `foo()` above is recursive. *It calls itself*. Let's try calling `foo()` and see what happens:

```
>>> from recursion import foo  
>>> foo(10)  
Traceback (most recent call last):  
  File "<stdin>", line 1, in <module>  
  File "./recursion.py", line 2, in foo  
    return foo(n-1)  
  File "./recursion.py", line 2, in foo  
    return foo(n-1)  
  File "./recursion.py", line 2, in foo  
    return foo(n-1)  
  File "./recursion.py", line 2, in foo  
    return foo(n-1)  
  File "./recursion.py", line 2, in foo  
    return foo(n-1)  
  File "./recursion.py", line 2, in foo  
    return foo(n-1) # etc. etc. etc.  
  File "./recursion.py", line 2, in foo  
    return foo(n-1)  
RuntimeError: maximum recursion depth exceeded
```

- Hmm. Our program crashed! What's going on? Well we initially invoke `foo(10)`, which invokes `foo(9)` which invokes `foo(8)` which invokes `foo(7)` which invokes `foo(6)` which invokes...
- Thus our initial `foo(10)` call is the first in an *infinite* sequence of calls to `foo()`. Computers do not like an infinite number of anything. For each of our `foo()` function invocations Python instantiates a data structure to represent that particular call to the function. That data structure is called a *stack frame*. A stack frame occupies memory. Our program attempts to create an infinite num-

ber of stack frames. That would require an infinite amount of memory. Our computer does not have an infinite amount of memory. So our program crashes (after a while).

- The problem with our recursive function is that it *never* fails to invoke itself and thus exhibits *infinite recursion*.
- To prevent infinite recursion we need to insert a *base case* into our function. Let's rewrite our function as `bar()` but this time cause it to stop once its parameter hits zero:

```
def bar(n):  
    if n == 0: # base case : no more calls to bar()  
        return 0  
    return bar(n-1)
```

- Let's try calling `bar()` and see what happens:

```
>>> from recursion import bar  
>>> bar(10)  
0
```

- Why does `bar` return zero? Well `bar(10)` calls `bar(9)` which calls `bar(8)` ... which calls `bar(0)`. The base case is `bar(0)`. It returns zero to `bar(1)` which returns zero to `bar(2)` which returns zero to `bar(3)` ... which returns zero to `bar(10)` which returns zero which is our answer.
- That's how recursion works. So far so good. But can we use recursion to do something useful?

Summing the numbers 0 through N

- Assume we have a function `sum_up_to()`. Given an argument `N` `sum_up_to(N)` returns the sum all of the integers 0 through `N`. For example `sum_up_to(10)` sums the sequence 10, 9, 8, 7, 6, 5, 4, 3, 2, 1, 0.
- Let's look at the `sum_up_to()` function in action:

```
>>> from recursion import sum_up_to  
>>> sum_up_to(10)  
55  
>>> list(range(11)) # let's verify we got the right answer  
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10]  
>>> sum(list(range(11)))  
55
```

- Let's try some more examples:

```
>>> from recursion import sum_up_to  
>>> sum_up_to(0)  
0  
>>> sum_up_to(1)  
1  
>>> sum_up_to(2)  
3  
>>> sum_up_to(3)  
6  
>>> sum_up_to(4)
```

```

10
>>> sum_up_to(5)
15
>>> sum_up_to(6)
21
>>> sum_up_to(7)
28
>>> sum_up_to(8)
36
>>> sum_up_to(9)
45
>>> sum_up_to(10)
55

```

- Do you notice anything recursive about the above sequence? Let's annotate each line to make the recursion obvious:

```

>>> from recursion import sum_up_to
>>> sum_up_to(0) # base case: returns zero
0
>>> sum_up_to(1) # returns 1 + sum_up_to(0)
1
>>> sum_up_to(2) # returns 2 + sum_up_to(1)
3
>>> sum_up_to(3) # returns 3 + sum_up_to(2)
6
>>> sum_up_to(4) # returns 4 + sum_up_to(3)
10
>>> sum_up_to(5) # returns 5 + sum_up_to(4)
15
>>> sum_up_to(6) # returns 6 + sum_up_to(5)
21
>>> sum_up_to(7) # returns 7 + sum_up_to(6)
28
>>> sum_up_to(8) # returns 8 + sum_up_to(7)
36
>>> sum_up_to(9) # returns 9 + sum_up_to(8)
45
>>> sum_up_to(10) # returns 10 + sum_up_to(9)
55

```

- For any argument N , $\text{sum_up_to}(N)$ is equal to $N + \text{sum_up_to}(N-1)$. This is the essence of a recursive solution. The solution to the problem $\text{sum_up_to}(N)$ is broken down into the operation $N +$ on a simpler version of the same problem $\text{sum_up_to}(N-1)$. For example $\text{sum_up_to}(10)$ is $10 + \text{sum_up_to}(9)$. The base case ensures recursion stops at some point. It encodes the fact that $\text{sum_up_to}(0)$ is zero.
- Let's write the Python code that implements the $\text{sum_up_to}()$ function:

```

def sum_up_to(n):
    if n == 0: # base case : no more calls to sum_up_to()
        return 0
    return n + sum_up_to(n-1)

```

- Why does $\text{sum_up_to}(10)$ return 55? Well $\text{sum_up_to}(10)$ calls $\text{sum_up_to}(9)$ which calls $\text{sum_up_to}(8)$... which calls $\text{sum_up_to}(0)$. The base case is $\text{sum_up_to}(0)$. It returns zero to $\text{sum_up_to}(1)$ which returns 1 ($1+0$) to $\text{sum_up_to}(2)$ which returns 3 ($2+1$) to $\text{sum_up_to}(3)$ which returns 6 ($3+3$) to $\text{sum_up_to}(4)$ which returns 10 ($4+6$) to $\text{sum_up_to}(5)$ which returns 15

(5+10) to `sum_up_to(6)` ... which returns 45 (9+36) to `sum_up_to(10)` which returns 55 (10+45) which is our answer.

Recursive factorial

- Factorial 4 or $4! = 4 * 3 * 2 * 1$ and in general $N! = N * (N-1) * (N-2) * (N-3) * \dots 2 * 1$.
- $1!$ is defined as 1.
- Let's look at some examples of factorial in action:

```
>>> from recursion import factorial
>>> factorial(1)
1
>>> factorial(2)
2
>>> factorial(3)
6
>>> factorial(4)
24
>>> factorial(5)
120
>>> factorial(6)
720
>>> factorial(7)
5040
>>> factorial(8)
40320
>>> factorial(9)
362880
>>> factorial(10)
3628800
```

- Do you notice anything recursive about the above sequence? Let's annotate each line to make the recursion obvious:

```
>>> from recursion import factorial
>>> factorial(1) # base case: returns 1
1
>>> factorial(2) # returns 2 * factorial(1)
2
>>> factorial(3) # returns 3 * factorial(2)
6
>>> factorial(4) # returns 4 * factorial(3)
24
>>> factorial(5) # returns 5 * factorial(4)
120
>>> factorial(6) # returns 6 * factorial(5)
720
>>> factorial(7) # returns 7 * factorial(6)
5040
>>> factorial(8) # returns 8 * factorial(7)
40320
>>> factorial(9) # returns 9 * factorial(8)
362880
>>> factorial(10) # returns 10 * factorial(9)
3628800
```

- For any argument N , `factorial(N)` is equal to $N * \text{factorial}(N-1)$. This is the essence of a recursive solution. The solution to the problem `factorial(N)` is broken down into the operation $N * \text{factorial}(N-1)$ on a simpler version of the same problem `factorial(N-1)`. For example `factorial(10)` is $10 * \text{factorial}(9)$. The base case ensures recursion stops at some point. It encodes the fact that `factorial(1)` is 1.
- Let's write the Python code that implements the `factorial()` function:

```
def factorial(n):
    if n == 1: # base case : no more calls to factorial()
        return 1
    return n * factorial(n-1)
```

Our old friend Fibonacci

- The Fibonacci sequence of numbers is given by: 1, 1, 2, 3, 5, 8, 13, etc. The first two numbers of the sequence are both defined to be 1 and thereafter each number in the sequence is defined as the sum of the previous two.
- Let's look at some examples of `fibonacci()` in action:

```
>>> from recursion import fibonacci
>>> fibonacci(0)
1
>>> fibonacci(1)
1
>>> fibonacci(2)
2
>>> fibonacci(3)
3
>>> fibonacci(4)
5
>>> fibonacci(5)
8
>>> fibonacci(6)
13
>>> fibonacci(7)
21
>>> fibonacci(8)
34
>>> fibonacci(9)
55
>>> fibonacci(10)
89
>>> fibonacci(11)
144
>>> fibonacci(12)
233
```

- Do you notice anything recursive about the above sequence? Let's annotate each line to make the recursion obvious:

```
>>> from recursion import fibonacci
>>> fibonacci(0) # base case: returns 1
1
>>> fibonacci(1) # base case: returns 1
1
```

```

>>> fibonacci(2) # fibonacci(1) + fibonacci(0)
2
>>> fibonacci(3) # fibonacci(2) + fibonacci(1)
3
>>> fibonacci(4) # fibonacci(3) + fibonacci(2)
5
>>> fibonacci(5) # fibonacci(4) + fibonacci(3)
8
>>> fibonacci(6) # fibonacci(5) + fibonacci(4)
13
>>> fibonacci(7) # fibonacci(6) + fibonacci(5)
21
>>> fibonacci(8) # fibonacci(7) + fibonacci(6)
34
>>> fibonacci(9) # fibonacci(8) + fibonacci(7)
55
>>> fibonacci(10) # fibonacci(9) + fibonacci(8)
89
>>> fibonacci(11) # fibonacci(10) + fibonacci(9)
144
>>> fibonacci(12) # fibonacci(11) + fibonacci(10)
233

```

- In general, $\text{fibonacci}(N) = \text{fibonacci}(N-1) + \text{fibonacci}(N-2)$. Our base cases are $\text{fibonacci}(0) = 1$ and $\text{fibonacci}(1) = 1$.
- Let's translate this into Python...

Reversing a list

- Let's try to come up with a recursive implementation of a function that reverses a list.
- Let's look at some examples of `reverse_list()` in action:

```

>>> from recursion import reverse_list
>>> reverse_list([])
[]
>>> reverse_list([5])
[5]
>>> reverse_list([4,5])
[5, 4]
>>> reverse_list([3,4,5])
[5, 4, 3]
>>> reverse_list([2,3,4,5])
[5, 4, 3, 2]
>>> reverse_list([1,2,3,4,5])
[5, 4, 3, 2, 1]

```

- Do you notice anything recursive about the above sequence? Let's annotate each line to make the recursion obvious:

```

>>> from recursion import reverse_list
>>> reverse_list([]) # base case : returns []
[]
>>> reverse_list([5]) # returns reverse_list([]).append(5)
[5]
>>> reverse_list([4,5]) # returns reverse_list([5]).append(4)
[5, 4]
>>> reverse_list([3,4,5]) # returns reverse_list([4,5]).append(3)
[5, 4, 3]

```

```
>>> reverse_list([2,3,4,5]) # returns reverse_list([3,4,5]).append(2)
[5, 4, 3, 2]
>>> reverse_list([1,2,3,4,5]) # returns reverse_list([2,3,4,5]).append(1)
[5, 4, 3, 2, 1]
```

- Let's translate this into Python...